

# Sparse Representations in Deep RL: Encoding Information in Deep Architectures

---

Hunter Hobbs

Anuj Pasricha

Matthew Resnick

# Problem

---

Unfortunately it hasn't changed:

## Catastrophic Interference



- A neural network containing a representation of a task can have significant overwriting of parameters when seeing new experiences.
- Could be a result of training on different regions of an environment for a single task or training on different tasks entirely.
- Five approaches to addressing this problem: regularization, ensembling, rehearsal, dual-memory, and sparse-coding (Kemker et. Al., 2017).
  - Our approaches are primarily sparse-coding.
  - A bit of regularization/ensembling here and there to make it happen.

# Theoretical Approach 1: Sparse Pathway Encoding

Idea: Train a deep controller normally, but encode input information onto the network structure itself.

We can do this with activation functions provided the power of trainable parameters:

- Maxout (Bengio et. Al., 2013)
- Channel-out (Wang & JaJa, 2013)



The learned network is not sparse, but a sparse representation exists as a subnetwork, accessible via activation.

# Theoretical Approach 2: Distributional Regularizers

First, **KL-divergence**: a measure of difference in information between two probability distributions. Think Shannon entropy.

Then, assume the output of each hidden unit is a random variable with some distribution  $p$ . We can use KL-divergence as a penalty term on the loss function to encourage  $p$  towards another distribution with a higher density about zero. Hence, a sparse representation!

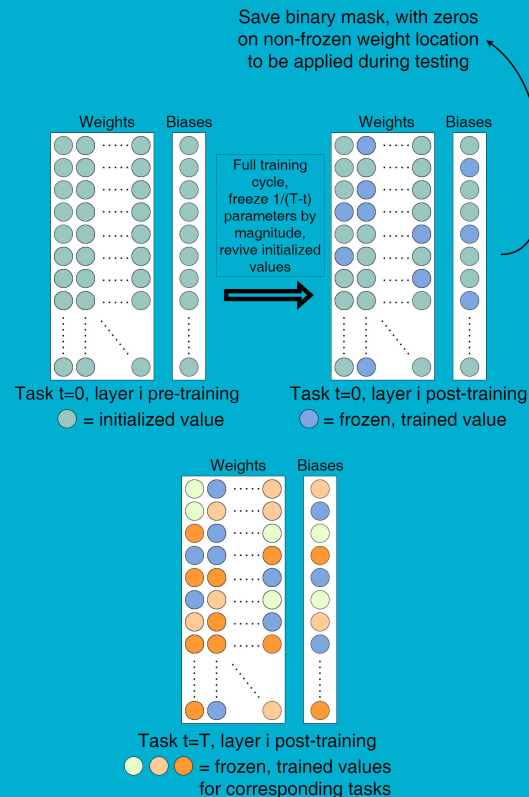
With a clipped KL-divergence, we can specify a minimum level of sparsity (Liu, 2019).

# Theoretical Approach 3: Spatially Packed Task Representations

Idea: We can prune many parameters from the network and still have a good representation of a task. Instead, we should just use the remaining space for more representations of more tasks. Same network size as before pruning, but we've learned multiple sparse representations.

For  $T$  total tasks, and task number  $t$ , we train normally, then freeze  $1/(T-t)$  tasks from training, and save a binary mask with 1s at these locations. Reinitialize, and repeat.

This is a modified PackNet, from Mallya et. Al. an algorithm for filters in a CNN for Computer Vision tasks.



**All approaches are modular, so can be applied to any Deep RL algorithm.**

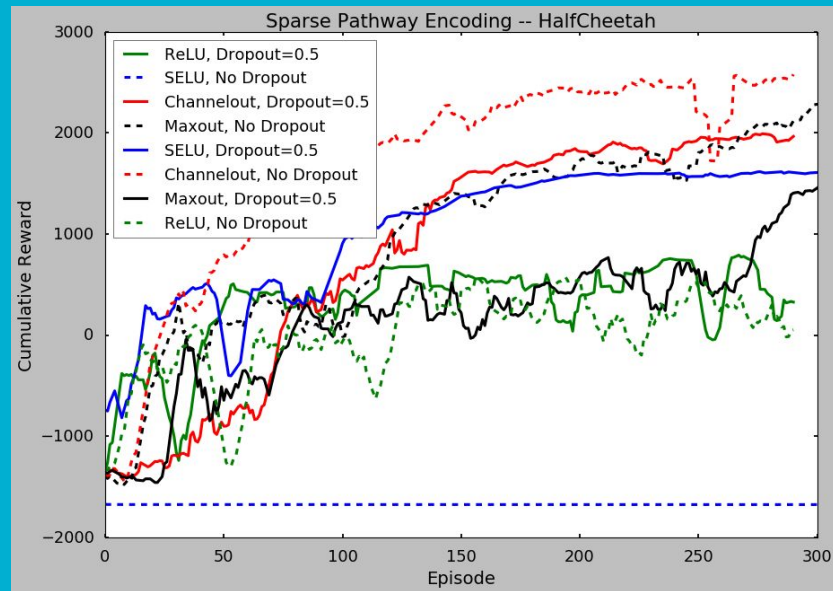

**The following results were obtained from applying the approaches to the actor network of a vanilla DDPG.**

**Info on other learning algorithms in the conclusion.**

# SPE Results

# Without PER


Training dynamics:





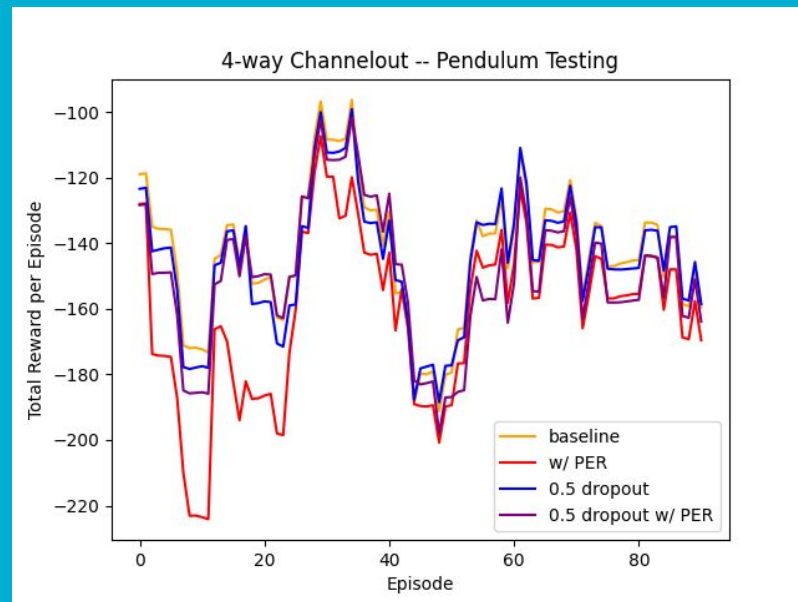

# Without PER

Testing:







# With PER

---



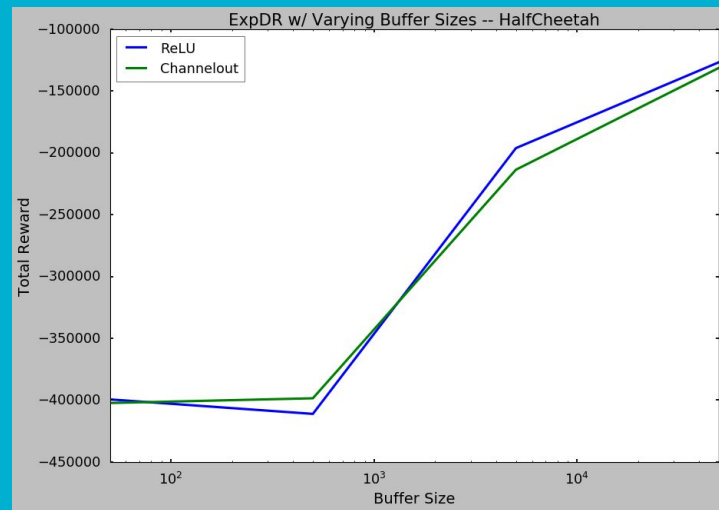
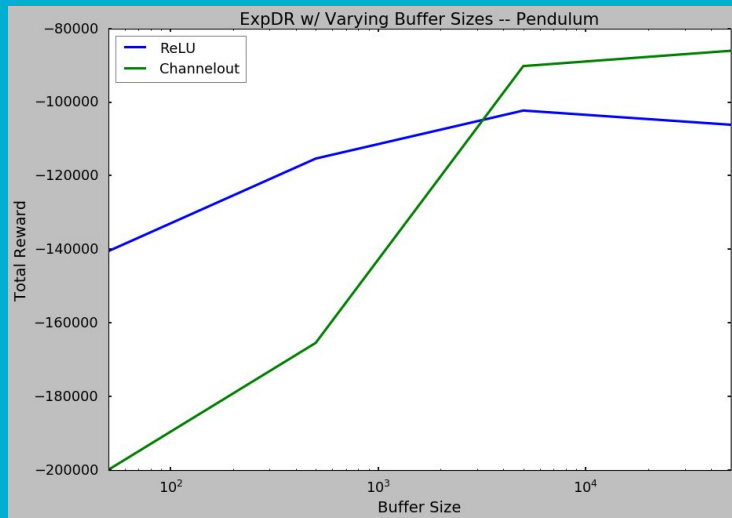
# Distributional Regularizers Results

# Without PER



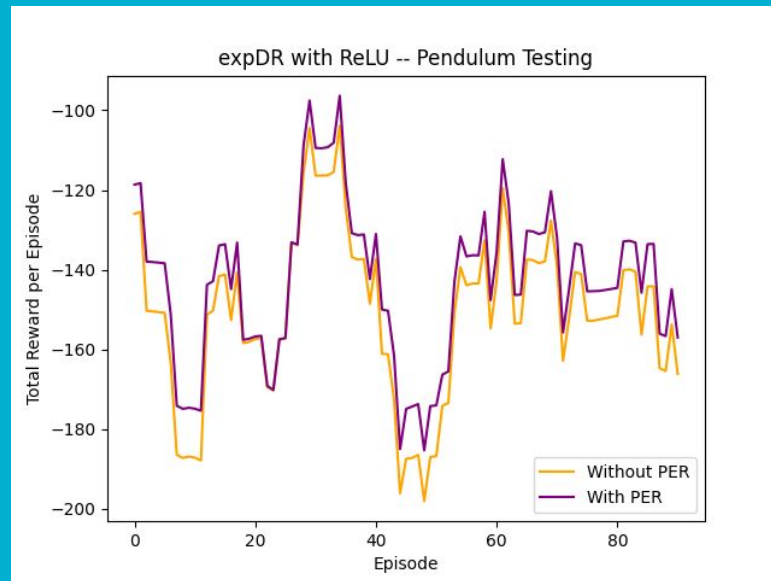
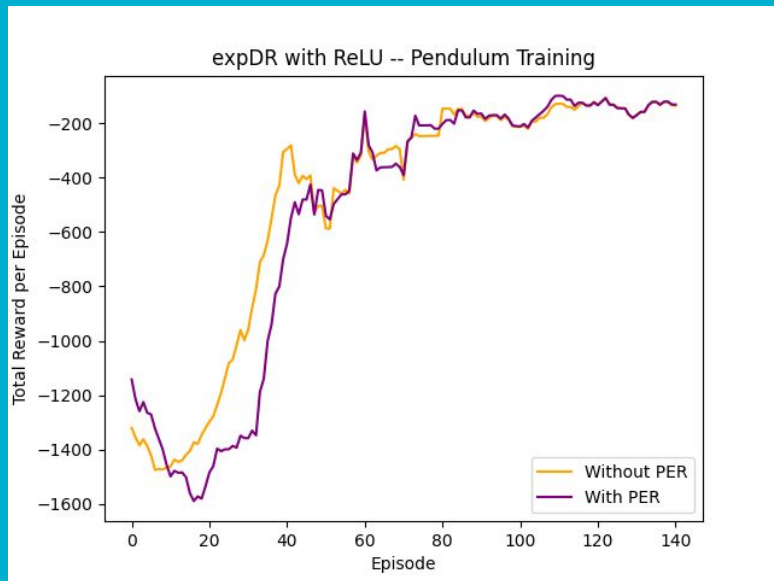
# Without PER

---




# With PER

---



# PackNet Results

Trained over three random seeds of Pendulum, 100 episodes each.



Model	Mean Reward
Channel-out Baseline	-351.41
Channel-out PackNet	-344.58
ReLU PackNet	-366.8
ReLU Baseline	-378.31

Testing over 100 episodes on the last seed and a completely new seed:


Model	Mean Reward
Channel-out Baseline (last)	-151.76
Channel-out Baseline (new)	-149.01
<b>Channel-out PackNet (last)</b>	<b>-135.91</b>
Channel-out PackNet (new)	-149.37
ReLU PackNet (last)	-151.03
ReLU PackNet (new)	-161.27
ReLU Baseline (last)	-157.56
ReLU Baseline (new)	-148.78

So PackNet, and especially with channel-out, works better than baseline. But these results are not very strong. More testing...




## Channel-out Activation

Trained over two random seeds of HalfCheetah,  
150 episodes each.



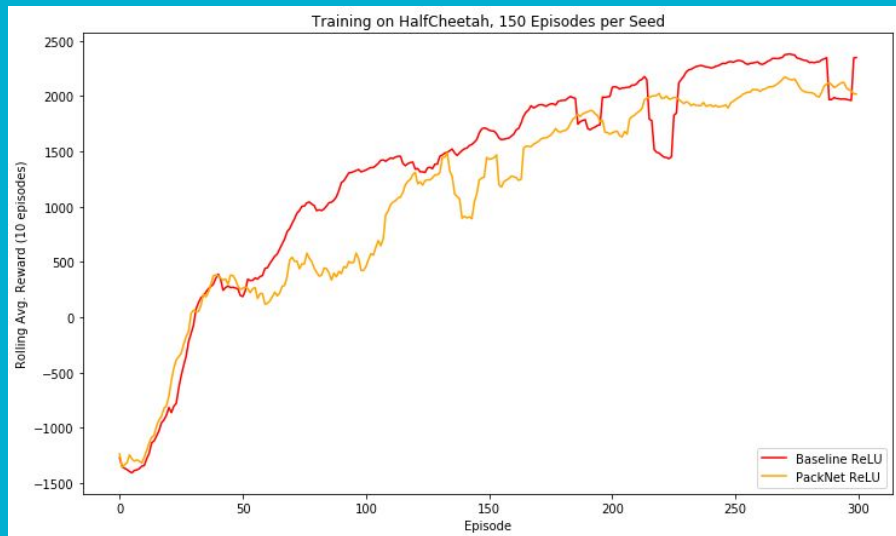
Testing on 100 episodes of 5 new random seeds.




Stronger results! But is it PackNet that causes this success, or its combination with channel-out?

## ReLU Activation

Trained over two random seeds of HalfCheetah,  
150 episodes each.



Testing on 100 episodes of 5 new random seeds.



So it's probably the combination. We claim that channel-out encodes observation information into the network's pathways, and PackNet increases network capacity usage like dropout, but without the added variance.

# Conclusion/Future Work

---

Main takeaways from Sparse Pathway Encoding:

Some activation function need dropout to be effective, but the variance effect of dropout is too great to make its use worthwhile with them.

SELU + alphaDropout shows it is possible to add a similar effect without using the vanilla notion of dropout.

# Conclusion/Future Work

---

Main takeaway from Distributional Regularizers:

ExpDR + ReLU is the only Distributional Regularizer model that worked well from a rewards standpoint, though all implementations were successful in creating a sparse representation.

- Our hypothesis about GaussianDR + channel-out was wrong! (Not too surprising because channel-out can't necessarily guarantee zero-output like ReLU, at the pool-level only)
- GaussianDR + ReLU worked well for DQN, what gives?
  - Is it the algorithms themselves, or do discrete vs continuous action spaces interact with sparse representations differently? Need more investigation to find out.

# Conclusion/Future Work

---

Main takeaways from PackNet:

PackNet and channel-out need each other.

- Channel-out captures observation information and encodes it into network pathways
- PackNet increases network capacity usage so the pathway information is more abundant and expressive.

PackNet can act as a limited alternative to dropout in Deep RL for some activation functions.

# References

---

- R. Kemker, M. McClure, A. Abitino, T. Hayes, and C. Kanan, “Measuring catastrophic forgetting in neural networks,” 2017.
- I. J. Goodfellow, D. Warde-Farley, M. Mirza, A. Courville, and Y. Bengio, “Maxout networks,” 2013.
- Q. Wang and J. JaJa, “From maxout to channel-out: Encoding information on sparse pathways,” 2013.
- V. Liu, “Sparse representation neural networks for online reinforcement learning,” Master’s thesis, University of Alberta, 2019.
- A. Mallya and S. Lazebnik, “Packnet: Adding multiple tasks to a single network by iterative pruning,” 2017.